Aradhya Chandra

Noida, India | aradhyachandra5@gmail.com | linkedin.com/in/aradhya-chandra-41390a30b | github.com/aradhyachandra

EDUCATION

Shiv Nadar Institution of Eminence, Greater Noida, India *2024 – 2028 (Expected)*
Bachelor of Technology in Computer Science & Engineering

TECHNICAL SKILLS

Languages: Python, SQL, Java, C, HTML/CSS
Data & ML: pandas, DuckDB, scikit-learn, LightGBM, time-series forecasting, RAG (LangChain, ChromaDB, Ollama)
Business Analytics: DCF and ratio analysis, churn and cohort analysis, inventory planning, scenario analysis
Tools: Git/GitHub, Streamlit, SQLite, Parquet, VS Code

EXPERIENCE

AI/ML Research Intern — IIT Kanpur *May 2026 – June 2026*
Summer Training Internship, under faculty guidance

- Built a domain-specific RAG-based legal assistant chatbot using Python, LangChain, ChromaDB, and Ollama (Llama 3.2), with a Streamlit front end.
- Designed the full retrieval pipeline end to end: PDF ingestion, recursive text chunking, persistent vector storage with nomic-embed-text embeddings, and distance-threshold-filtered retrieval to reduce hallucination.

Marketing Associate — Finvest, SNU's Finance Club *May 2025 – Present*

- Selected for the marketing team to promote the university's annual finance and business conclave; coordinated with a 10+ member team on on-ground strategies that increased footfall by 20%.

PROJECTS

Retail Demand Forecasting & Inventory Planning (Walmart M5) | Python, SQL (DuckDB), LightGBM, Streamlit *Oct 2026*

- Forecast 28-day daily demand for 5,748 Walmart grocery item-stores (11.2M rows), prepared in SQL with DuckDB; found most items sell only intermittently (86% of series), which shaped the baselines and metrics.
- Benchmarked a global LightGBM model against 7 baselines (incl. Croston and TSB) in a 4-window rolling backtest, then froze it and ran the unseen 28-day holdout once: **about 5% more accurate than the best simple rule and 24% more accurate than same-weekday-last-week**, matching the backtest.
- Fed the forecasts into a reorder policy with safety stock and simulated it on actual demand: **5.5% less inventory than the best rule and 10.4% less than the naive rule for the same 95% fill rate**, holding across all 20 cost and lead-time scenarios.
- Spotted that the dataset's prices are derived from sales (changes in 9 of 221,569 no-sale weeks) and redesigned price features to avoid leakage; showed 10% of item-stores carry 93% of lost-sales cost.

AI Equity Research Assistant (TCS) | Python, SQLite, ChromaDB, Ollama, Streamlit *Oct 2026*

- Built a research assistant where a local LLM picks tools from fixed lists and writes the answer, while Python computes every figure (11 ratios, DCF) from 178 facts in three TCS annual reports, each checked against its page; a verifier rejects unsupported numbers.
- On a frozen held-out test against a standard RAG baseline using the same model, scored **15 of 17 numeric answers correct vs 2 of 17**, with **0 of 40 unsourced numeric claims vs 4 of 18**.
- Fed a DCF engine from the verified filings and checked every assumption against four years of history; the check flagged two 7% growth years above TCS's best recent year, so I cut the base case to a flat 6%.

Customer Churn Prediction & Retention Analytics | Python, pandas, scikit-learn, Streamlit *Oct 2026*

- Predicted which customers of a UK online retailer would stop buying within 90 days from 1.07M real transactions (5,875 customers), using 21 behavioural features built from leakage-proof time snapshots.
- Tuned on an earlier month, froze the model, and tested it once on a later month: gradient boosting caught **83% true churners in its top 10% of risk scores, vs 74% for a recency rule and 50.5% by chance**.
- Turned the scores into a costed retention plan across 27 margin, save-rate and offer-cost scenarios; found the campaign's value depended more on the assumptions than on the model, and recommended a small controlled pilot to measure the real save rate before rollout.

CERTIFICATIONS

- Summer Training Program – Python for Data Science, IIT Kanpur (E&ICT Academy) — Sep 2025
- Penn Engineering: Introduction to Python Programming (May 2024); Introduction to Java and Object-Oriented Programming (May 2026)