Aradhya Chandra

Noida, India | aradhyachandra5@gmail.com | linkedin.com/in/aradhya-chandra | github.com/aradhyachandra

EDUCATION

Shiv Nadar Institution of Eminence, Greater Noida, India *2024 – 2028 (Expected)*
Bachelor of Technology in Computer Science & Engineering

TECHNICAL SKILLS

Languages: Python, SQL, Java, C, HTML/CSS
Data & ML: pandas, DuckDB, scikit-learn, LightGBM, time-series forecasting, RAG (LangChain, ChromaDB, Ollama)
Business Analytics: DCF and ratio analysis, churn and cohort analysis, inventory planning, scenario analysis
Tools: Git/GitHub, Streamlit, SQLite, Parquet, VS Code

EXPERIENCE

AI/ML Research Intern — IIT Kanpur *May 2026 – June 2026*
Summer Training Internship, under faculty guidance

- Built a domain-specific RAG-based legal assistant chatbot using Python, LangChain, ChromaDB, and Ollama (Llama 3.2), with a Streamlit front end.
- Designed the full retrieval pipeline end to end: PDF ingestion, recursive text chunking, persistent vector storage with nomic-embed-text embeddings, and distance-threshold-filtered retrieval to reduce hallucination.

Marketing Associate — Finvest, SNU's Finance Club *May 2025 – Present*

- Selected for the marketing team to promote the university's annual finance and business conclave; coordinated with a 10+ member team on on-ground strategies that increased footfall by 20%.

PROJECTS

Retail Demand Forecasting & Inventory Planning (Walmart M5) | Python, SQL (DuckDB), LightGBM, Streamlit *Oct 2026*

- Forecast 28-day daily demand for 5,748 Walmart grocery item-stores (11.2M rows), prepared in SQL with DuckDB; found most items sell only intermittently (86% of series), which shaped the baselines and metrics.
- Benchmarked a global LightGBM model against 7 baselines (incl. Croston and TSB) in a 4-window rolling backtest, then froze it and ran the unseen 28-day holdout once: **about 5% more accurate than the best simple rule and 24% more accurate than same-weekday-last-week**, matching the backtest.
- Fed the forecasts into a reorder policy with safety stock and simulated it on actual demand: **5.5% less inventory than the best rule and 10.4% less than the naive rule for the same 95% fill rate**, holding across all 20 cost and lead-time scenarios.
- Spotted that the dataset's prices are derived from sales (changes in 9 of 221,569 no-sale weeks) and redesigned price features to avoid leakage; showed 10% of item-stores carry 93% of lost-sales cost.

AI Equity Research Assistant (TCS) | Python, SQLite, ChromaDB, Ollama, Streamlit *Oct 2026*

- Built a research assistant where a local LLM only routes and explains, while Python tools compute every figure (11 ratios, DCF) from 178 hand-verified facts in three TCS annual reports; a verifier blocks any number no tool produced.
- On a frozen held-out test against a standard RAG baseline using the same model, scored **15 of 17 numeric answers correct vs 2 of 17**, with **0 of 40 unsourced numeric claims vs 4 of 18**.
- Fed a DCF engine from the verified filings and checked every assumption against four years of history automatically, which led to revising base-case revenue growth from 7% to 6%.

Customer Churn Prediction & Retention Analytics | Python, pandas, scikit-learn, Streamlit *Oct 2026*

- Predicted which customers of a UK online retailer would stop buying within 90 days from 1.07M real transactions (5,875 customers), using 21 behavioural features built from leakage-proof time snapshots.
- Tuned on an earlier month, froze the model, and tested it once on a later month: gradient boosting caught **83% true churners in its top 10% of risk scores, vs 74% for a recency rule and 50.5% by chance**.
- Turned the scores into a costed retention plan across 27 margin, save-rate and offer-cost scenarios; found the campaign's value depended more on the assumptions than on the model, and recommended a small controlled pilot to measure the real save rate before rollout.

CERTIFICATIONS

- Summer Training Program – Python for Data Science, IIT Kanpur (E&ICT Academy) — Sep 2025
- Penn Engineering: Introduction to Python Programming (May 2024); Introduction to Java and Object-Oriented Programming (May 2026)